

Parallel Finite Element assembly

Michael Wathen*

May 6, 2016

1 Introduction

Many industrial and geophysical scientific computing problems require discrete solving techniques for partial differential equations (PDEs). The two key components of a PDE solve are:

- discretisation;
- linear/non-linear solve.

For discretisation there are three main methods which are used: finite differences, finite volumes and finite element methods (FEM). FEM discretises the model and represents the solution with respect to basis functions. Such methods rely heavily on variational methods from calculus of variations to approximate the solution.

The second key component is, the solving technique. Often, for large three-dimensional scientific computing problems it is not possible to do a direct solve for the system matrix. This is because such problems still have a large band width (lots of entries far away from the diagonal), and hence the memory and time consumption is too large. Therefore, an iterative approach is required for these problems.

For a efficient FEM implementation the work to construct an $n \times n$ matrix is $\mathcal{O}(n)$. However, this is order is not necessarily true for all PDE solve. For example, if the viscosity is small then it is known that the Navier-Stokes equations become harder to solve. There is a huge amount of work that goes into the development for optimal ($\mathcal{O}(n)$) solvers for PDE; for example see [3, 4, 6, 10].

*Department of Computer Science, The University of British Columbia, Vancouver, BC, V6T 1Z4, Canada, mwathen@cs.ubc.ca.

For this project we will mainly focus on parallel finite element method within the **FEniCS** software framework [8]. We will show assembly times for both the Laplacian and magnetohydrodynamics equations. Since the Laplacian plays an integral roll in many PDEs, we will also look at the timings for a parallel solve of it.

2 FEM in FEniCS

To understand the basic principle of parallel finite element methods, it is necessary to look at the sequential version. Given a certain cell, \mathcal{T}_h , with the local-to-global degrees of freedom mapping i_T then the generic finite element assembly algorithm is given by:

Algorithm 1 FEM assembly

```

1:  $A = 0$ 
2: for  $\mathcal{T} \in \mathcal{T}_h$  do
3:   (1) Compute  $i_T$ 
4:   (2) Compute  $A_T$ 
5:   (3) Add  $A_T$  to  $A$  according to  $i_T$ :
6:   for  $i \in \mathcal{I}_T$  do
7:      $A_{i_{T(i)}} \stackrel{+}{=} A_{T,i}$ 

```

end

From this algorithm, we see that the general principle of the FEM is to loop through the cells of the discretised domain and calculate the local cell matrix entries. Then the local elements are mapped back into the global matrix, A . This therefore means that for certain global degrees of freedom require multiple reads and writes.

The algorithm for a basic parallel finite element method does not really vary from the sequential version but assembles over a partitioned mesh. Therefore, for an efficient parallel implementation an appropriate mesh partitioning is required. The partitioned mesh should distribute the minimises the inter-process communication costs between the degrees of freedom. **FEniCS** does not do the mesh partitioning by it self, however it does have backends for the two most prominent mesh partitioning software packages **METIS** [7] and **SCOTCH** [9].

3 Results

In this section we will provide two examples. First, we consider a simple Laplacian where we look at both the assembly and solve times in parallel. Second, we look at the parallel assembly time for the magnetohydrodynamics problem.

The two numerical experiment have been carried out using the finite element software **FEniCS** [8] together with **PETSc4PY** package (Python interface for **PETSc** [1, 2]) and the multigrid package **HYPRE** [5]. The mesh partitioning software used was **SCOTCH**.

3.1 Laplacian

For our first example, consider Laplace’s equation with non-homogeneous Dirichlet boundary conditions.

$$\begin{aligned}\Delta u &= f \text{ in } \Omega, \\ u &= g \text{ on } \partial\Omega.\end{aligned}\tag{1}$$

The Laplacian appears in fluid dynamics, electromagnetism and many other applications. Therefore, there has been a significant amount of work that has gone into efficient solvers for this system. It is well known that multigrid (either geometric (GM) or algebraic multigrid (AMG)) yields an optimal, $\mathcal{O}(n)$, solving algorithm.

The assembly and solve results are shown in Tables 1 and 2, respectively. First, we look at the scalability with respect to the problem size. For the

MPI processes	DoFs				
	14,739	107,811	823,875	6,440,067	50,923,779
1	6.44e-01	2.69e+00	1.86e+01	1.49e+02	-
2	1.91e-01	1.38e+00	1.04e+01	7.92e+01	-
4	1.29e-01	1.05e+00	5.66e+00	4.19e+01	-
8	8.71e-02	5.31e-01	3.11e+00	2.32e+01	1.88e+02
16	1.16e-01	4.34e-01	2.12e+00	1.34e+01	9.94e+01
32	2.48e-01	3.17e-01	1.69e+00	1.20e+01	9.14e+01

Table 1: Laplacian assembly time for different degrees of freedom and MPI processes

set up considered, the problems increase by a factor of 8 per level. From the tables, we observe that the that as we increase the degrees of freedom

MPI processes	DoFs				
	14,739	107,811	823,875	6,440,067	50,923,779
1	3.67e-01	4.50e+00	4.34e+01	3.93e+02	-
2	2.12e-01	2.90e+00	2.79e+01	1.92e+02	-
4	1.74e-01	1.46e+00	1.40e+01	1.21e+02	-
8	8.42e-02	1.18e+00	1.18e+01	9.24e+01	7.76e+02
16	8.28e-02	1.26e+00	9.11e+00	8.00e+01	6.71e+02
32	6.27e-02	9.35e-01	8.70e+00	7.66e+01	6.50e+02

Table 2: Laplacian solve time for different degrees of freedom and MPI processes

then both the assembly and solve time increase by roughly a factor of 8 for a fixed number of MPI processes.

When varying the number of MPI processes the results don't scale quite as well, at least for the solve time. From Table 1 the times do seem to decrease by a factor of 2 when the number of MPI processes are increased from 1 to 16. However, moving 32 processes does not decrease the time by a factor of 2 from 16. The solve times do not exhibit this behaviour, though the timings do decrease as the number of MPI processes increase.

3.2 MHD

For our second example, we will consider an incompressible magnetohydrodynamics model with Dirichlet boundary conditions:

$$-\nu \Delta \mathbf{u} + (\mathbf{u} \cdot \nabla) \mathbf{u} + \nabla p - \kappa (\nabla \times \mathbf{b}) \times \mathbf{b} = \mathbf{f} \quad \text{in } \Omega, \quad (2a)$$

$$\nabla \cdot \mathbf{u} = 0 \quad \text{in } \Omega, \quad (2b)$$

$$\kappa \nu_m \nabla \times (\nabla \times \mathbf{b}) + \nabla r - \kappa \nabla \times (\mathbf{u} \times \mathbf{b}) = \mathbf{g} \quad \text{in } \Omega, \quad (2c)$$

$$\nabla \cdot \mathbf{b} = 0 \quad \text{in } \Omega. \quad (2d)$$

In a standard FEM fashion, we linearize around the current velocity and magnetic fields and introduce basis functions corresponding to the discrete spaces as in [11]. This yields the following matrix system:

$$\begin{pmatrix} F(u) & B^T & C(b)^T & 0 \\ B & 0 & 0 & 0 \\ -C(b) & 0 & M & D^T \\ 0 & 0 & D & 0 \end{pmatrix} \begin{pmatrix} \delta u \\ \delta p \\ \delta b \\ \delta r \end{pmatrix} = \begin{pmatrix} r_u \\ r_p \\ r_b \\ r_r \end{pmatrix}, \quad (3)$$

with

$$\begin{aligned} r_u &= f - F(u)u - C(b)^T b - B^T p, \\ r_p &= -Bu, \\ r_b &= g - Mu + C(b)b - D^T r, \\ r_r &= -Db, \end{aligned}$$

where $F(u) = A + O(u)$. The matrices are: F , the discrete convection-diffusion operator; B , a discrete divergence operator; M , the discrete curl-curl operator; D , a discrete divergence operator; and C , a discrete coupling term. For full discretisation details see [12, Chapter 2].

At this point there is no known parallel implementation of this discretisation for the MHD model with these specific elements used in [11]. In the recent paper [10], the authors look at an exact penalty form of the MHD model. This discretisation yields a 3×3 linear system to be solved at each non-linear iteration. Their implementation appears to be the first parallel approach for of both the solve and finite element discretisation.

The finite element assembly results are presented in Table 3. Due to the

MPI processes	DoFs				
	20,381	148,661	1,134,437	8,861,381	70,045,061
1	4.82e+00	3.94e+01	3.04e+02	2.46e+03	-
2	3.08e+00	2.09e+01	1.58e+02	1.30e+03	-
4	1.50e+00	1.06e+01	8.49e+01	6.60e+02	-
8	7.88e-01	5.86e+00	4.65e+01	3.58e+02	2.81e+03
16	6.72e-01	3.33e+00	2.40e+01	1.88e+02	1.47e+03
32	6.66e-01	3.29e+00	2.35e+01	1.85e+02	2.01e+03

Table 3: MHD assembly time for different degrees of freedom and MPI processes

much greater complexity of the MHD model (2) than the Laplacian (1), we would expect that the system assembly times to be generally higher for a similar number problem size. This is exactly what we observe between the Tables 3 and 1. As with the Laplacian assembly results the MHD model does seem to scale with the problem size. Again, we see a similar scaling behaviour with respect to the number of MPI processes until we reach 32 MPI threads. We see that the computation time stalls or possible increases from 16 to 32 processes.

4 Conclusion

We have looked at a parallel finite element discretisation of the well-known Laplacian as well as new MPI implementation of the MHD model described in [11]. For our numerical experiments we have used the comprehensive FEniCS finite element software package.

The Laplacian is a well understood problem, but it forms an integral part of many PDEs. Therefore, it is a good starting point for numerical test. We saw that the parallel assembly times for the Laplacian exhibit the optimal ($\mathcal{O}(n)$) times across specified MPI processors. This is particularly important for the larger three-dimensional meshes. We also presented parallel solve times for the Laplacian. Again, the solver exhibited optimal, $\mathcal{O}(n)$, times with respect to the mesh.

Finally, we looked at the more complex MHD model. MHD models electrically conductive fluids in the presence of an electromagnetic field. So far there are very few parallel finite element implementations of these models. The results shown in Section 3.2 show good scalability with respect to the mesh. We see that it takes roughly the same time to construct a matrix of dimension 8,861,381 on one processor as it does to construct a problem of 70,045,061 on eight processors. This is only the starting point of parallel MHD. The aim is to implement a parallel solve using block based preconditioners similar to [10, 12] to create a fully parallel and scalable solution method.

References

- [1] S. Balay, M. F. Adams, J. Brown, P. Brune, K. Buschelman, V. Eijkhout, W. D. Gropp, D. Kaushik, M. G. Knepley, L. C. McInnes, K. Rupp, B. F. Smith, and H. Zhang. PETSc User’s Manual. Technical Report ANL-95/11 - Revision 3.4, Argonne National Laboratory, 2013.
- [2] S. Balay, M. F. Adams, J. Brown, P. Brune, K. Buschelman, V. Eijkhout, W. D. Gropp, D. Kaushik, M. G. Knepley, L. C. McInnes, K. Rupp, B. F. Smith, and H. Zhang. PETSc Web page. <http://www.mcs.anl.gov/petsc>, 2014.
- [3] J. Bosch, d. Kay, m. Stoll, and A. Wathen. Fast solvers for cahn–hilliard inpainting. *SIAM Journal on Imaging Sciences*, 7(1):67–97, 2014.

- [4] H. C. Elman, D. J. Silvester, and A. J. Wathen. *Finite Elements and Fast Iterative Solvers: with Applications in Incompressible Fluid Dynamics*. Oxford University Press, second edition, 2014.
- [5] R. D. Falgout and U. Yang. *hypr*: A library of high performance preconditioners. In *Computational Science — ICCS 2002*, volume 2331 of *Lecture Notes in Computer Science*, pages 632–641. Springer Berlin Heidelberg, 2002.
- [6] C. Greif and D. Schötzau. Preconditioners for the discretized time-harmonic Maxwell equations in mixed form. *Numerical Linear Algebra with Applications*, 14(4):281–297, 2007.
- [7] G. Karypis. METIS: Family of multilevel partitioning algorithms, 2010. <http://glaros.dtc.umn.edu/>.
- [8] A. Logg, K. A. Mardal, and G. N. Wells, editors. *Automated solution of differential equations by the finite element method*, volume 84 of *Lecture Notes in Computational Science and Engineering*. Springer, Heidelberg, 2012.
- [9] F. Pellegrini. Scotch. <http://www.labri.fr/perso/pelegrin/scotch>.
- [10] E. G. Phillips, H. C. Elman, E. C. Cyr, J. N. Shadid, and R. P. Pawlowski. A block preconditioner for an exact penalty formulation for stationary MHD. *SIAM Journal on Scientific Computing*, 36(6):B930–B951, 2014.
- [11] D. Schötzau. Mixed finite element methods for stationary incompressible magneto–hydrodynamics. *Numerische Mathematik*, 96(4):771–800, 2004.
- [12] M. Wathen. Iterative Solution of a Mixed Finite Element Discretisation of an Incompressible Magnetohydrodynamics Problem. Master’s thesis, University of British Columbia, 2014.